Summary of Neural Network Models

Project of machine learning

Comparative Analysis of Neural Network Models for Credit Default Prediction

Contents

Abstract

This report presents a comparative analysis between different neural network architectures aimed at predicting credit defaults. Utilizing a real-world dataset, various neural network models are implemented and evaluated based on accuracy, computational efficiency, and reproducibility.

Introduction

The prediction of credit defaults is a critical task for financial institutions. Neural networks have been recognized as an effective method for this purpose. This project examines various neural network architectures to identify the most efficient model in terms of predictive performance and computational requirements.

The dataset is obtained from Kaggle and the existing codes on the platform are not referenced.

Methodology

Data Preparation

The dataset[1] was processed using Pandas and preprocessed for neural network application. Categorical features are transformed by one hot encoding method[2]. Standardization of scalar features was performed using Standard Scaler from Scikit-Learn to maintain fairness in model comparisons. The dataset was subsequently divided into training and testing sets, adhering to an 80-20 split.

Neural Network Models

Several neural network models were implemented utilizing TensorFlow. A base model comprising two dense layers with ReLU activation and a sigmoid activation output layer was established to set a benchmark for subsequent comparisons.

Experimental Setup

Training of the models occurred on the designated training set, incorporating a 20% validation split. The Adam optimizer and binary cross-entropy loss were uniformly applied across all models to maintain consistency in the experimental conditions. Hyperparameters were selected based on prior experiments and literature reviews to mitigate overfitting and underfitting.

1. **Model Compilation[3]:**

---

[1] https://www.kaggle.com/datasets/uciml/default-of-credit-card-clients-dataset/data. This dataset contains information on default payments, demographic factors, credit data, history of payment, and bill statements of credit card clients in Taiwan from April 2005 to September 2005. Variable list is in appendix.

[2] 'SEX', 'EDUCATION', 'MARRIAGE', 'PAY_0', 'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6'. An interesting finding is that, without transformation of categorical features, the accuracy level does not vary significantly.

[3] Base model:

```
# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
# Train the model
original_history = model.fit(X_train, y_train, epochs=20, batch_size=32, validation_split=0.2)
```

- **Input Layer:** defined by the input shape (X_train.shape[1],) which corresponds to the number of features in the input data.
- **Hidden Layers:** Two hidden layers are defined with 128 and 64 neurons, respectively, both using the ReLU (Rectified Linear Unit) activation function. ReLU is a commonly used activation function that introduces non-linearity into the model.
- **Output Layer:** The output layer consists of a single neuron with a sigmoid activation function, which is suitable for binary classification tasks. Sigmoid function squashes the output to a range between 0 and 1, representing the probability of the input belonging to the positive class (class 1).
- **Optimizer**: 'adam'. Adam is an adaptive learning rate optimization algorithm that's well-suited for training deep neural networks. It computes adaptive learning rates for each parameter.
- **Loss Function**: 'binary_crossentropy'. This is a common choice for binary classification problems. It measures the difference between the true labels and the predicted probabilities.
- **Metrics**: ['accuracy']. During training, accuracy will be monitored to evaluate the performance of the model.

2. **Model Training**:

- **Epochs**: Number of times the entire dataset is passed forward and backward through the neural network.
- **Batch Size**: Number of samples propagated through the network before the weights are updated. Smaller batch sizes tend to provide a regularizing effect and can lead to faster convergence.

3. **Model Evaluation**:

- The model is evaluated on the test set using the previously specified loss function and accuracy metric.

Theory Background:

In binary classification problems, the output layer typically consists of a single neuron with a sigmoid activation function, which produces values between 0 and 1, representing the probability of the input belonging to one of the two classes[4].

Math Equations:

1. **ReLU Activation Function**:

$$f(x) = \max(0, x)$$

- ReLU is defined as the positive part of its argument x. It introduces non-linearity by outputting the input directly if it is positive, otherwise, it outputs zero.
- **Avoid vanishing gradient problem(not like sigmoid)**

---

[4] One softmax variation of model was tested in the last step to check the activation function change could bring any change to the model. All other steps I used only sigmoid.

- o **Computational simplicity**
- o **Sparse activation(negative neurons are not activated)**
- o **Improved convergence**

2. **Sigmoid Activation Function**: $\sigma(z) = \frac{1}{1+e^{-z}}$

   - o Used in the output layer to squash the output between 0 and 1, representing the probability of the input belonging to class 1. **Which is better for the model objective.**
   - o **Smooth gradient. This function is differentiable, suitable for gradient based optimization methods and backpropagation.**
   - o **Interpretable as the probability**
   - o **Limitations: vanishing gradient, computational expensive, non zero centered output.**

3. **Binary Crossentropy Loss Function**:

$$L(y, \hat{y}) = -\frac{1}{N}\sum_{i=1}^{N} [y_i \log(\hat{y}_i) + (1 - y_i)\log(1 - \hat{y}_i)]$$

   - o Where y represents the true labels (0 or 1), $\hat{y}$ represents the predicted probabilities, and N is the number of samples.
   - o **It is suited for sigmoid output function to constrain output range**
   - o **Penalizing incorrect classifications more heavily when model is more confident on wrong predictions, considered as strong gradient signal**
   - o **Encouraging probability extremes**
   - o **Gradient friendly**

4. **Softmax:**

$$\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

   - o **Multi class classification problem**
   - o **convert a vector of arbitrary real values into a vector of probabilities**

5. **Sparse Categorical Cross-Entropy:**

$$L(y, \hat{y}) = -\sum_i y_i \log(\hat{y}_i)$$

   - o **Suitable to pair with softmax**
   - o It computes the cross-entropy loss between the true labels and the predicted probability distribution.

6. **Adam Optimization Algorithm**:

- o Adam updates the weights using the gradient of the loss function with respect to the weights. It adapts the learning rates of each weight parameter by estimating the first and second moments of the gradients.

Hyperparameters:

- **Structure**:
  - o Input Layer: Determined implicitly by the shape of the input data.
  - o Output Layer: Single neuron with sigmoid activation for binary classification.
  - o Hidden Layers: Number of hidden layers and neurons per layer are not specified in this code snippet.

- **Other**:
  - o Activation Functions: ReLU for hidden layers, sigmoid for the output layer.
  - o Optimizer: adam
  - o Loss Function: binary_crossentropy
  - o Metrics: accuracy
  - o Epochs
  - o Batch Size

Regularization:

1. **L1 Regularization (Lasso Regression):**

   L1 regularization adds a penalty term to the loss function equal to the absolute value of the weights.Mathematically, the loss function with L1 regularization can be expressed as:

$$\text{Loss}_{\text{L1}} = \text{Loss}_{\text{original}} + \lambda \sum_{i=1}^{n} |w_i|$$

where:

   - o $\text{Loss}_{\text{original}}$ is the original loss function (e.g., cross-entropy loss or mean squared error).
   - o λ is the regularization parameter that controls the strength of regularization.
   - o w_i are the weights of the neural network.
   - o The regularization term $\lambda \sum_{i=1}^{n} |w_i|$ encourages sparsity in the weight matrix, as it tends to shrink less important weights to zero.

2. **L2 Regularization (Ridge Regression):**

   L2 regularization adds a penalty term to the loss function equal to the squared magnitude of the weights. Mathematically, the loss function with L2 regularization can be expressed as:

$$\text{Loss}_{\text{L2}} = \text{Loss}_{\text{original}} + \lambda \sum_{i=1}^{n} w_i^2$$

where:

- $\text{Loss}_{\text{original}}$ is the original loss function.
- λ is the regularization parameter.
- w_i are the weights of the neural network.
- The regularization term $\lambda \sum_{i=1}^{n} w_i^2$ penalizes large weights, encouraging them to be small but does not force them to zero.

In practice, during training, the optimizer minimizes the regularized loss function (original loss function plus the regularization term) with respect to the weights. The choice between L1 and L2 regularization depends on the specific problem and the desired properties of the learned model. L1 regularization tends to produce sparse solutions, while L2 regularization tends to produce more diffuse solutions with smaller weights. Sometimes, a combination of both (elastic net regularization) is used to leverage the benefits of both methods.

Neural network model training and hyperparameter tuning

Steps are followed as training, repeat(try additional three times to see the change of metrics), adjustment, fine-tuning(with different hyperparameter groups).

1. **model 0 and re-running(they are the same model runed for 4 times):**

Model	Training Accuracy (%)	Validation Accuracy (%)	Training Loss(%)	Validation Loss(%)	Test Accuracy (%)
0	**83.06**	**80.04**	**39.82**	**46.42**	**81.45**
1	83.42	80.62	39.39	45.71	81.02
2	83.23	80.75	39.56	45.23	81.40
3	83.11	81.31	39.82	44.95	81.17

Table 1 Four independent executions of training same nerual networks
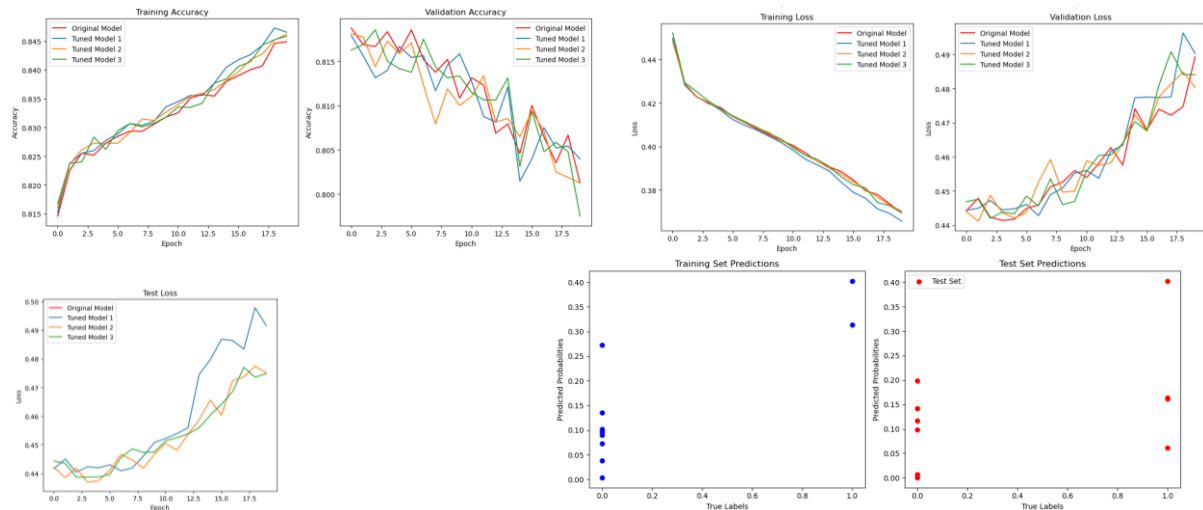
Figure 1 Reruning of base model, training, validation, test accuracy and loss, and sampled dataset test on accuracy

Observation:

1. Each training session provides different result, although the differences of the ranges are not evident.
2. By random sampling test, the predicted values are evidently different from the real values of the label.

Problems detected: overfitting.

2. Then, tuning the hyperparameters[5]

Variation	Train Accuracy	Validation Accuracy	Test Accuracy	Train Loss	Validation Loss	Test Loss
Var 1	84.724%	79.813%	80.350%	0.363166	0.492602	0.490032
Var 2	84.177%	80.292%	81.200%	0.373027	0.479086	0.476731
Var 3	84.412%	79.896%	80.600%	0.371068	0.480192	0.477003
Var 4	84.646%	80.646%	81.383%	0.368240	0.491831	0.481148
Var 5	85.012%	79.875%	80.267%	0.362796	0.494384	0.488361
Var 6	84.546%	79.583%	81.083%	0.370867	0.485538	0.473637
Var 7	84.724%	80.167%	80.267%	0.367261	0.480633	0.480172
Var 8	84.896%	79.792%	80.850%	0.368417	0.493958	0.487366
Var 9	84.486%	79.542%	80.933%	0.371354	0.517088	0.474893
Var 10	84.583%	80.125%	80.950%	0.368786	0.484298	0.469284

Table 2 tuning hyperparameters

Observations:

1. Relatively small variation of model performances
2. No significant improvement compared to previous results(though every time the code was run the results turn different.

Problems detected: overfitting. No huge difference between model variations.

---

[5] creat more models to fine tune : epochs +- 10 holding others constant. Batch +-10 holding others constant. Validation split+- 0,1 holding others constant.
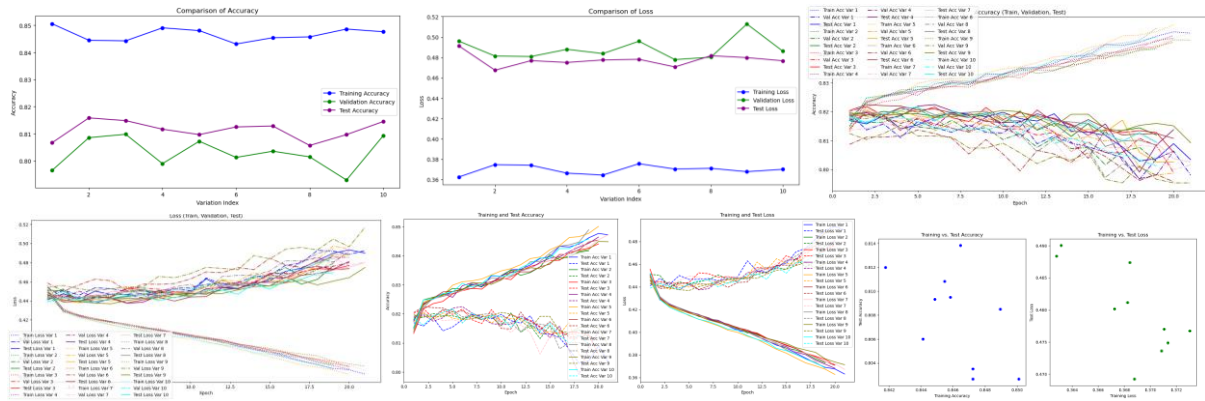
Figure 2 parameter tunning models comparison, on accuracy, loss, epoch-wise training, validation, test accuracy and losses, then training and testing accuracy and loss comparison, then training and test accuracy and loss scatter plot

3. Last, tuning[6] neuron number hyperparameters[7]

Model Configuration	Test Accuracy (%)	Training Accuracy (%)	Validation Accuracy (%)	Training Loss	Validation Loss
64 neurons first layer, 32 neurons second	81.95	82.78	81.02	0.4044	0.4518
64 neurons first layer, 48 neurons second	81.30	82.86	81.10	0.4058	0.4487
64 neurons first layer, 64 neurons second	81.53	82.79	81.48	0.4028	0.4420
64 neurons first layer, 80 neurons second	81.32	83.05	81.12	0.3996	0.4526
64 neurons first layer, 96 neurons second	81.28	83.17	81.31	0.4009	0.4497
64 neurons first layer, 112 neurons second	80.78	83.21	80.52	0.3984	0.4638
96 neurons first layer, 32 neurons second	81.52	82.82	81.08	0.4026	0.4478
96 neurons first layer, 48 neurons second	**82.02**	**84.26**	**80.87**	**0.3765**	**0.4762**
96 neurons first layer, 64 neurons second	81.30	83.20	80.81	0.3969	0.4536
96 neurons first layer, 80 neurons second	81.35	83.26	81.12	0.3964	0.4510
96 neurons first layer, 96 neurons second	81.25	83.22	81.00	0.3946	0.4497
96 neurons first layer, 112 neurons second	81.45	83.05	81.10	0.3951	0.4542
128 neurons first layer, 32 neurons second	81.68	83.02	81.46	0.4024	0.4482
128 neurons first layer, 48 neurons second	81.63	83.11	81.04	0.3993	0.4550
128 neurons first layer, 64 neurons second	81.23	83.22	81.27	0.3964	0.4569
128 neurons first layer, 80 neurons second	81.17	83.21	81.02	0.3961	0.4504
128 neurons first layer, 96 neurons second	81.10	83.32	81.38	0.3950	0.4542
128 neurons first layer, 112 neurons second	81.68	83.46	81.00	0.3907	0.4563
160 neurons first layer, 32 neurons second	81.45	82.85	81.10	0.3978	0.4595
160 neurons first layer, 48 neurons second	81.32	83.39	81.17	0.3946	0.4450

---

[6] Additional sample data finetuning is not possible. All fintuning steps are focusing on adjustments of hyperparameters/structure of the model.

[7] All results and metrics for the models change everytime the codes are run.

Model Configuration	Test Accuracy (%)	Training Accuracy (%)	Validation Accuracy (%)	Training Loss	Validation Loss
160 neurons first layer, 64 neurons second	81.20	83.22	80.69	0.3954	0.4527
160 neurons first layer, 80 neurons second	81.35	83.37	81.25	0.3942	0.4532
160 neurons first layer, 96 neurons second	81.37	83.21	80.96	0.3922	0.4592
160 neurons first layer, 112 neurons second	80.85	83.56	80.65	0.3850	0.4589
192 neurons first layer, 32 neurons second	81.40	83.01	81.21	0.3995	0.4496
192 neurons first layer, 48 neurons second	80.48	83.36	80.52	0.3960	0.4652
192 neurons first layer, 64 neurons second	81.43	83.33	81.35	0.3926	0.4602
192 neurons first layer, 80 neurons second	81.47	83.38	81.10	0.3895	0.4591
192 neurons first layer, 96 neurons second	81.52	83.51	80.87	0.3890	0.4548
192 neurons first layer, 112 neurons second	80.45	83.67	80.40	0.3878	0.4631
224 neurons first layer, 32 neurons second	81.50	83.23	81.17	0.3965	0.4552
224 neurons first layer, 48 neurons second	80.98	83.14	80.69	0.3954	0.4674
224 neurons first layer, 64 neurons second	81.58	83.27	81.23	0.3946	0.4533
224 neurons first layer, 80 neurons second	81.60	83.18	81.19	0.3918	0.4566
224 neurons first layer, 96 neurons second	81.25	83.49	80.81	0.3882	0.4544
224 neurons first layer, 112 neurons second	80.97	83.52	80.63	0.3866	0.4565

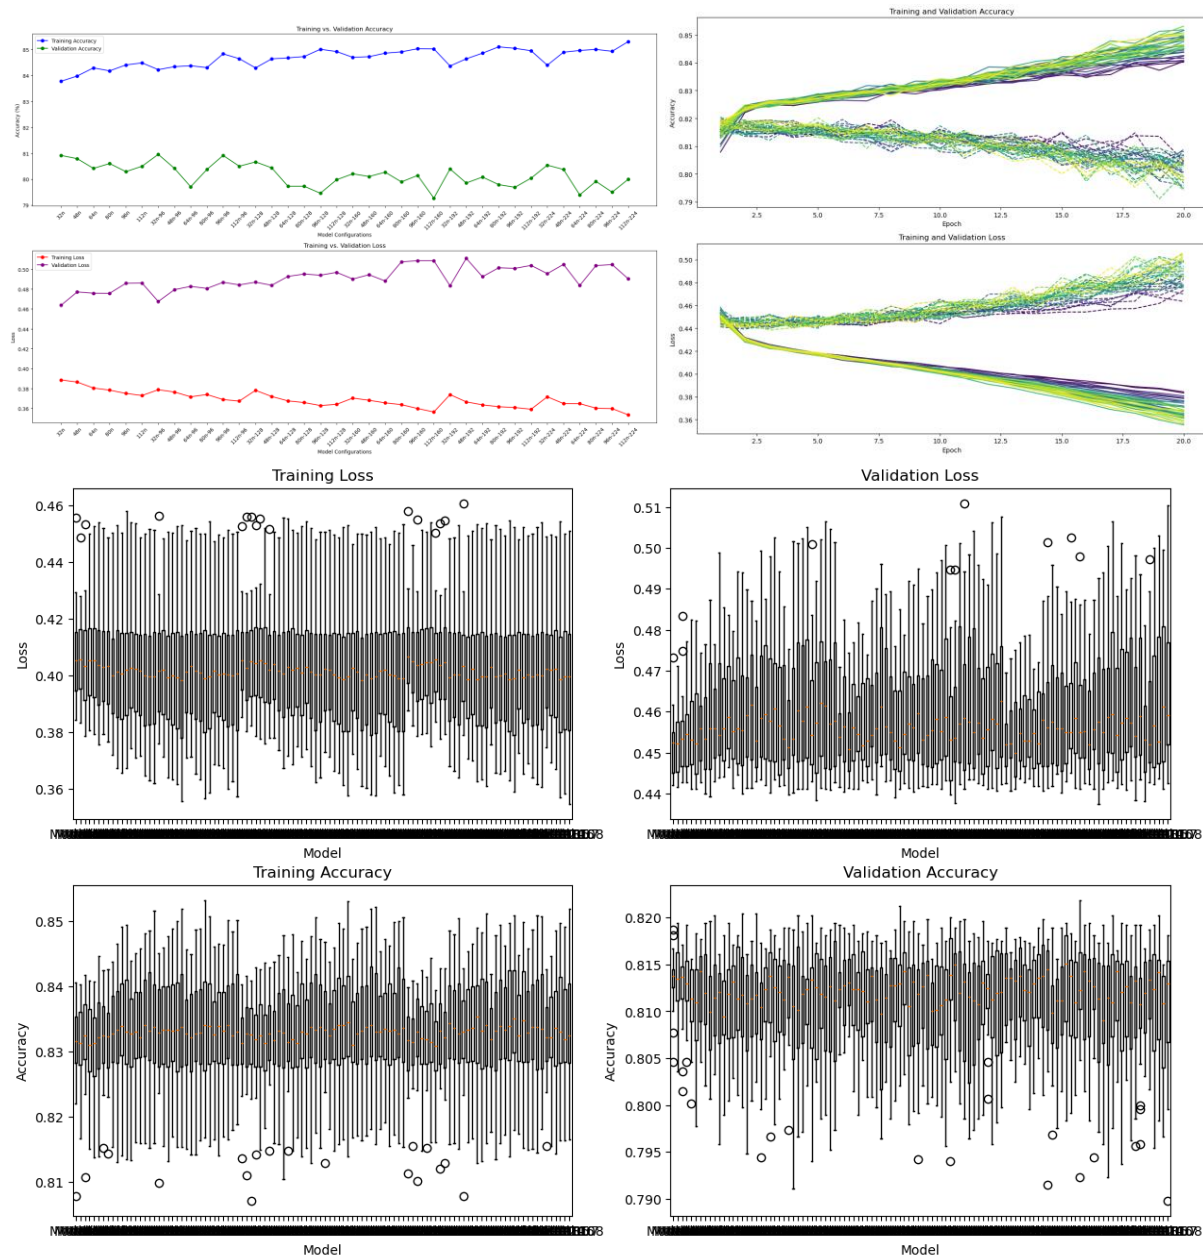*Table 3 fintuning hyperparameters on neuron numbers in each layer*

Figure 3 Hyperparameter Tuned models accuracy, loss comparison on training, validation, on static base, epoch wise, and the box plot to show fluctuation ranges of loss, accuracy on training validation

Observations:

Problems detected: overfitting. No huge difference between model variations.

Resolving the overfitting problem

Best Model Configuration: 96 neurons (1st layer), 48 neurons (2nd layer)

Based on the above model, the combination of L1 L2 regularization strategies are set up to find the best regularization strategy and its solution to improve the model.

The it is found that the best Model by Metrics is Number 4 in the below table.

	Test Accuracy	Training Accuracy	Validation Accuracy	Training Loss	Validation Loss	Regularization Type	L1 Strength	L2 Strength
0	0.820000	0.822708	0.818542	0.441258	0.452167	L1	0.001	0
1	0.817833	0.820781	0.814375	0.477329	0.484645	L1	0.010	0
2	0.781167	0.778490	0.777083	0.688834	0.690012	L1	0.100	0
3	0.819167	0.826979	0.816667	0.425356	0.444518	L2	0.000	0.001
4	**0.820667**	**0.822917**	**0.815208**	**0.439274**	**0.446761**	**L2**	**0.000**	**0.01**
5	0.811167	0.819635	0.808333	0.456318	0.460663	L2	0.000	0.1
6	0.819667	0.822031	0.816667	0.450602	0.457831	L1_L2	0.001	0.001
7	0.781167	0.778490	0.777083	0.546248	0.548758	L1_L2	0.001	0.01
8	0.781167	0.778490	0.777083	0.694016	0.696148	L1_L2	0.001	0.1
9	0.781167	0.778490	0.777083	0.546851	0.548813	L1_L2	0.010	0.001
10	0.781167	0.778490	0.777083	0.561146	0.562477	L1_L2	0.010	0.01
11	0.781167	0.778490	0.777083	0.707333	0.709565	L1_L2	0.010	0.1
12	0.781167	0.778490	0.777083	0.693895	0.694178	L1_L2	0.100	0.001
13	0.781167	0.778490	0.777083	0.708385	0.708829	L1_L2	0.100	0.01
14	0.781167	0.778490	0.777083	0.856683	0.860662	L1_L2	0.100	0.1

Table 4 L1 L2 combinations to regularize the best model from thre previous step to solve overfitting problem
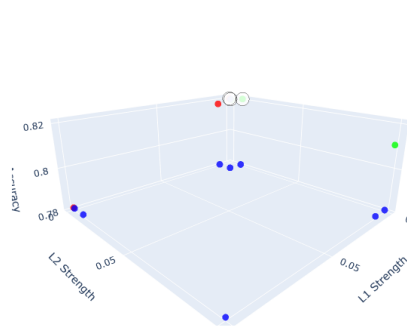

Figure 4 Scatterplot matrix about different regularization types, in 3D
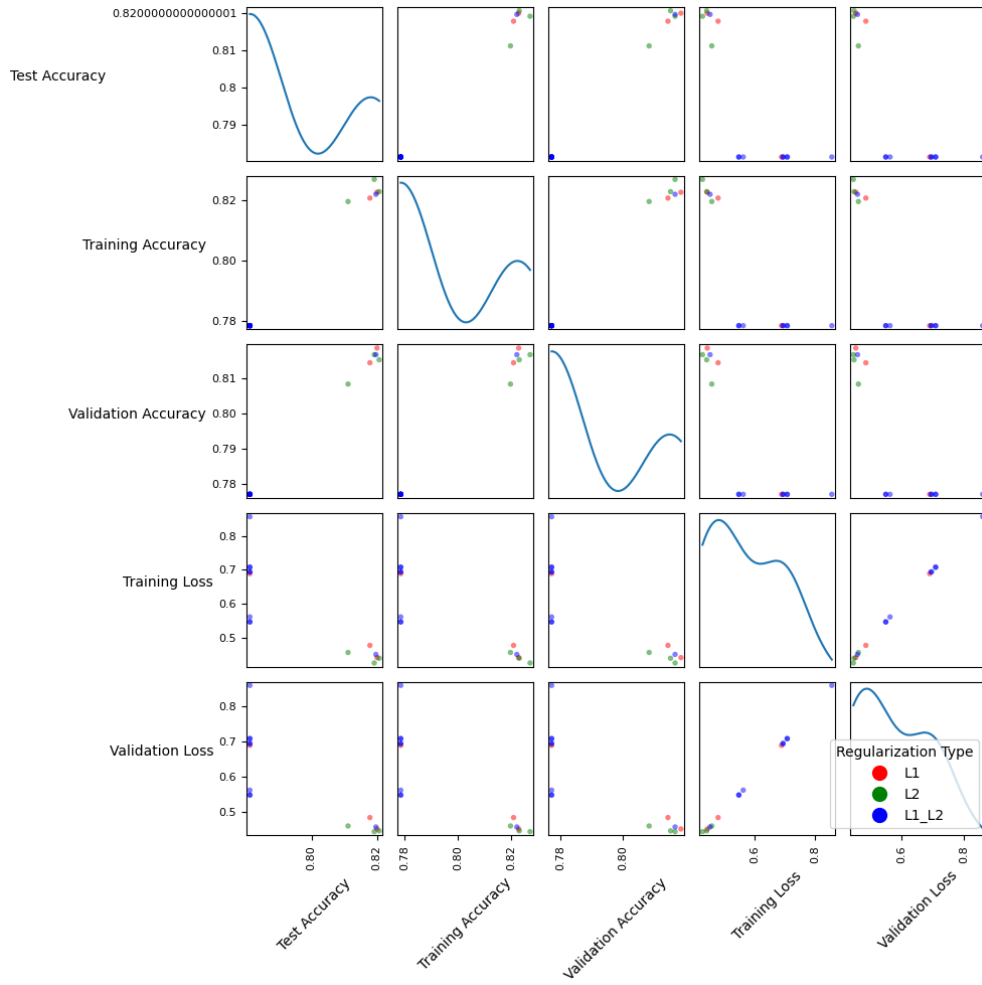
Figure 5 scatter plot of train and test accuracy of regularization choices, and train/validation accuracy and loss scatter plots


Figure 6 Comparison on the model used for activation function tunning, left as softmax function version and right as sigmoid function version[8]

---

[8] Softmax activation + Categorical Cross-Entropy loss for multi-class classification. Sigmoid activation + Binary Cross-Entropy loss for binary classification.

Figure 7 ORC comparison, left softmax model, right sigmoid model

From above results, Sigmoid choice makes the model training accuracy converges in the end.

For the sigmoid model:

- The training loss is 0.4346, the training accuracy is 82.35%.
- The validation loss is 0.4448, the validation accuracy is 81.69%.
- The test loss is 0.4401, the test accuracy is 82.15%.

For the softmax model:

- The training loss is 0.4341, the training accuracy is 82.41%.
- The validation loss is 0.4474, the validation accuracy is 81.25%.
- The test loss is 0.4434, the test accuracy is 81.92%.

Both models have similar performance metrics across all datasets, indicating that in this scenario, the choice of activation function (sigmoid or softmax) doesn't significantly impact the model's performance.

Results

Model performance was primarily evaluated using accuracy metrics on the testing dataset. The best models(from each section of test and training)' test accuracy, training and validation loss trend, etc, provide a benchmark for subsequent model comparisons.

After identifying best model on the hyperparameter tuning step, further regularization test was done to find next best model choice. Eventually a better model was produced. And, the model is very sensitive to the tunning strength change.

The reproducibility of results is also examined and maintained, by providing clear coding structure, emphasizing the transparent and precise documentation of experimental setups and parameters[9].

---

[9] Details in the code documents.

Additional models comparison

Other models

Previously trained ML models, including logistic, random forest, KNN, are listed below with metrics:

Model[10]	Accuracy	Precision	Recall	F1-Score	ROC AUC
Logistic Regression	0.8193	0.7617	0.2390	0.3639	0.7189
Random Forest	0.8213	0.6459	0.3840	0.4816	0.7655
K-Nearest Neighbors	0.7893	0.5190	0.3470	0.4159	0.7101
Logistic Ridge	0.809	0.680119	0.232143	0.346139	0.723051

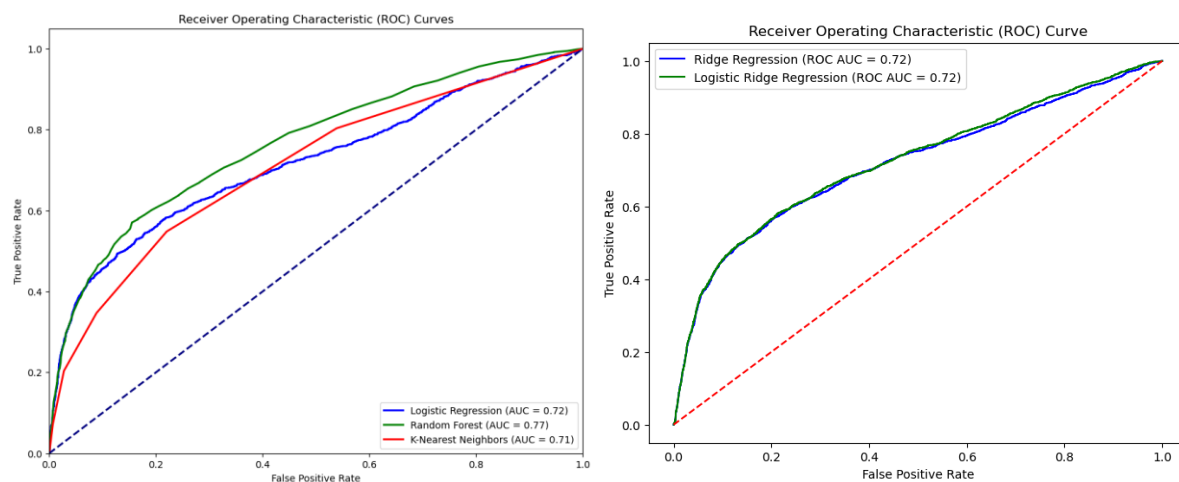*Table 5 Model metrics comparisons*

Figure 8 Additional algorithms comparison and logistic vs logistic ridge models

Below are the runtime records[11]:

Figure 9 Runtime statistics of logistic(similar to ridge logistic regression), random forest, KNN models, and, Runtime detail bar chart of Neural Network codes

---

[10] All models' codes are attached in github.

[11] Details in appendix.

Introducing Neural Network algorithms, the single step run time record and collective comparisons with all models(only on training session) are as follows[12]:

Figure 10 Comparison of all models regarding single training session runtime

Logistic Regression(and Ridge logistic regression[13]**):** Fastest training time at only 0.1 seconds, suitable for simpler datasets and problems.

Random Forest: More time-consuming at 4.7 seconds due to its complex nature of training multiple decision trees.

k-Nearest Neighbors (k-NN): Slower still at 11.2 seconds, particularly for large datasets due to the need for calculating distances between instances.

Single Neural Network Model: Takes the longest time to train at 20.4 seconds, reflecting the complexity and depth involved in neural network training.

Training neural network models has the highest computational resource usage and the finetuning results may not be higher than random forest models, but the above data has shown the result that any neural network model has outperformed logistic regression model[14] and ridge logistic regression model[15].

Conclusion

The investigation underscores the capability of neural networks in the context of credit default prediction. The findings indicate that the model accuracy level is among the highest(KNN without addressing overfitting problem shows accuracy similar or superior) with other models previously tested, considering additional efforts to tune the hyperparameters, and introduce regularization techniques to avoid overfitting problem.

Although the interpretation of neural network model is not straight forward like other simpler models, and the runtime level costs computational resources, it could be a viable choice to model credit default data set. Further research could investigate more sophisticated architectures, parameter tunning, and feature addition/engineering strategies to enhance predictive accuracy.

There are multiple improvements which could be done in the future analysis.

---

[12] After each execution of the codes, the results vary. Small variance can be ignored.
[13] Ridge logistic regression runtime level is similar, close to 0.
[14] This does not exclude the possibility where, logistic models can be improved.
[15] Accuracy: 0.809. But the penalty variation does not improve the model. Code is attached in github.

1. Changing hyper parameters by adding additional hidden layers in the model, or more aggressive adjustment.
2. Additional sample data could be validated and used as input to further verify or finetune the model.
3. Models should serve the use case. Considerations could include, structure descriptions and result interpretation, computational resources, time cost[16], other model usage constraints, etc.

Appendices

Supplementary materials encompass the project's code, detailing data preprocessing, model construction, training protocols, and evaluation metrics, ensuring the reproducibility of the experiments.

Variable list

Variable	Description
ID	ID of each client
LIMIT_BAL	Amount of given credit in NT dollars (includes individual and family/supplementary credit)
SEX	Gender (1=male, 2=female)
EDUCATION	Education level (1=graduate school, 2=university, 3=high school, 4=others, 5=unknown, 6=unknown)
MARRIAGE	Marital status (1=married, 2=single, 3=others)
AGE	Age in years
PAY_0	Repayment status in September, 2005 (-1=pay duly, 1=payment delay for one month, …, 9=payment delay for nine months and above)
PAY_2	Repayment status in August, 2005 (scale same as above)
PAY_3	Repayment status in July, 2005 (scale same as above)
PAY_4	Repayment status in June, 2005 (scale same as above)
PAY_5	Repayment status in May, 2005 (scale same as above)
PAY_6	Repayment status in April, 2005 (scale same as above)
BILL_AMT1	Amount of bill statement in September, 2005 (NT dollar)
BILL_AMT2	Amount of bill statement in August, 2005 (NT dollar)
BILL_AMT3	Amount of bill statement in July, 2005 (NT dollar)
BILL_AMT4	Amount of bill statement in June, 2005 (NT dollar)
BILL_AMT5	Amount of bill statement in May, 2005 (NT dollar)
BILL_AMT6	Amount of bill statement in April, 2005 (NT dollar)

[16] The whole project takes me around 2 months time excluding non working days. Where simple models took around 2-3 weeks and neural network model took more than one month. The time spent consists of model choosing, design, coding, fintune, result interpretations.

Variable	Description
PAY_AMT1	Amount of previous payment in September, 2005 (NT dollar)
PAY_AMT2	Amount of previous payment in August, 2005 (NT dollar)
PAY_AMT3	Amount of previous payment in July, 2005 (NT dollar)
PAY_AMT4	Amount of previous payment in June, 2005 (NT dollar)
PAY_AMT5	Amount of previous payment in May, 2005 (NT dollar)
PAY_AMT6	Amount of previous payment in April, 2005 (NT dollar)
default.payment.next.month	Default payment (1=yes, 0=no) for this dataset

Table 6 variable list of the dataset

Code[17]

As shown in github[18] and attached in the email.

Runtime details

Three single algorithms time breakdown:

Data Loading: Took 2.2 seconds. This process involves reading the file from disk, which is typically limited by the I/O speed of the system rather than the CPU.

Logistic Regression Training: Took only 0.1 seconds. Logistic Regression is faster compared to other models as it involves solving a convex optimization problem that usually converges quickly.

Random Forest Training: Took 4.7 seconds. This is more time-consuming due to the need to train multiple decision trees, which increases computational and memory requirements.

k-Nearest Neighbors (k-NN) Training: Took the longest at 11.2 seconds. k-NN can be slow, especially with large datasets, as it requires computing the distance to every instance in the training set.

All in One Code + Graph: Took 13.3 seconds. This includes the cumulative time of training all models, making predictions, and generating ROC curves for each.

Neural Network Development Time Breakdown:

Section 1: Single Model

Data Preprocessing: 3.5s

Model Architecture: 0s

Training: 20.4s

Evaluation & Additional Training: 54.5s

Visual Comparison: 1.1s

Sampling Cases: 1.8s

Output Model Layers: 0s

Section 2: Hyperparameter Finetuning

Variations Setting (10 models): 3m 16.4s

---

[17] Results in this report and in the markdown area of the code are not finite. Each execution of the code produces different results.

[18] https://github.com/StevenFromUnimiMIEDSE/Unimi-DSE-Machine-Learning-Project

Plot Function Preparation: 0s

Print Summaries and Charts: 1s

Section 3: Model Structure Finetuning

Changing Number of Neurons (36 models): 11m 7.8s

Print Summaries, Metrics, Graphs: 16.1s

Compare Models Collectively: 0.7s

Box Plot for Metric Variations: 1.5s

Shaded Variance Trend Graph: 0.6s

Best Model Selection: 0s